



Async. programming with ranges

Corehard Autumn 2018

dr Ivan Čukić

ivan@cukic.co
<http://cukic.co>



core-dev @ KDE
www.kde.org

About me

- KDE development
- Talks and teaching
- Author of "Functional Programming in C++" (Manning)
- Functional programming enthusiast, but not a purist

Disclaimer


Make your code readable. Pretend the next person who looks at your code is a psychopath and they know where you live.

Philip Wadler

ITERATORS

Iterators

At the core of generic programming in STL.




Iterators

```
std::vector<std::string> identifiers;  
  
std::copy_if(  
    std::cbegin(tokens),  
    std::cend(tokens),  
    std::back_inserter(identifiers),  
    is_valid_identifier);
```



Iterators

```
std::vector<std::string> identifiers;  
  
std::copy_if(  
    std::istream_iterator<std::string>(std::cin),  
    std::istream_iterator<std::string>(),  
    std::back_inserter(identifiers),  
    is_valid_identifier);
```


Example



Example



Example



Example



Example

```
template <typename It, typename Pred>
std::pair<It, It> move_selected(
    It begin, It end, It destination,
    Pred is_selected)
{
    return {
        std::stable_partition(begin, destination,
                               std::not_fn(is_selected)),
        std::stable_partition(destination, end,
                               is_selected)
    };
}
```

RANGES

Problems with iterators

Example: Sum of squares

This should be separable into two subtasks:

- squaring the values
- summation

Problems with iterators

```
std::vector<int> xs { ... };

std::vector<int> squared(xs.size());
std::transform(std::cbegin(xs), std::cend(xs),
               std::begin(squared),
               [](int x) { return x * x; });

return std::accumulate(std::cbegin(squared),
                       std::cend(squared),
                       0);
```

Problems with iterators

```
std::vector<int> xs { ... };  
  
return std::transform_reduce(par,  
    std::cbegin(xs), std::cend(xs),  
    std::cbegin(xs),  
    0);
```

Proxy iterators

```
template <typename Trafo, typename It>
class transform_proxy_iterator {
public:
    auto operator*() const
    {
        return transformation(
                *real_iterator);
    }

private:
    Trafo transformation;
    It real_iterator;
};
```

Ranges

```
[ iterator, sentinel )
```

Iterator:

- `*i` – access the value
- `++i` – move to the next element

Sentinel:

- `i == s` – has iterator reached the end?

Ranges

```
std::vector<int> xs { ... };
```

```
accumulate(  
    transform(  
        xs,  
        [] (int x) { ... },  
        0);
```


Ranges

```
std::vector<int> xs { ... };  
  
accumulate(  
    transform(  
        filter(  
            xs,  
            [] (int x) { ... } ),  
            [] (int x) { ... } ),  
    0);
```


Ranges

```
std::vector<int> xs { ... };
```

```
accumulate(  
    xs | filter([] (int x) { ... })  
      | transform([] (int x) { ... } ),  
    0 );
```

Word frequency

1986: Donald Knuth was asked to implement a program for the "Programming pearls" column in the Communications of ACM journal.

The task: Read a file of text, determine the n most frequently used words, and print out a sorted list of those words along with their frequencies.

Word frequency

1986: Donald Knuth was asked to implement a program for the "Programming pearls" column in the Communications of ACM journal.

The task: Read a file of text, determine the *n* most frequently used words, and print out a sorted list of those words along with their frequencies.


His solution written in Pascal was **10** pages long.

Word frequency

Response by Doug McIlroy was a 6-line shell script that did the same:

```
tr -cs A-Za-z '\n' |  
  tr A-Z a-z |  
  sort |  
  uniq -c |  
  sort -rn |  
  sed ${1}q
```


Functional thinking – data transformation



With ranges

```
const auto words =  
    // Range of words  
    istream_range<std::string>(std::cin)  
  
    // Cleaning the input  
    | view::transform(string_to_lower)  
    | view::transform(string_only_alnum)  
    | view::remove_if(&std::string::empty)  
  
    // Sort the result  
    | to_vector | action::sort;
```


With ranges

```
const auto results =  
    words  
  
    // Group repeated words  
    | view::group_by(std::equal_to<>())  
  
    // Count how many repetitions we have  
    | view::transform([] (const auto &group) {  
        return std::make_pair(  
            distance(begin(group), end(group),  
            *begin(group));  
        })  
  
    // Sort by frequency  
    | to_vector | action::sort;
```

With ranges

```
for (auto value:  
    results | view::reverse // descending  
           | view::take(n) // we need first n results  
    ) {  
    std::cout << value.first << " "  
               << value.second << std::endl;  
}
```

PUSH

Ranges

```
[ iterator, sentinel )
```

Iterator:

- $*i$ – access the value
- $++i$ – move to the next element

Sentinel:

- $i == s$ – has iterator reached the end

Ranges

```
[ iterator, sentinel )
```

Iterator:

- `*i` – access the value
- `++i` – move to the next element

BLOCKING!

Sentinel:

- `i == s` – has iterator reached the end

Push iterators



Push iterators

Each *push iterator* can:

- Accept values
- Emit values

No need for the accepted and emitted values to be 1-to-1.

Push iterators

- Sources – push iterators that only emit values
- Sinks – push iterators that only accept values
- Modifiers – push iterators that both accept and emit values

Continuation

```
template <typename Cont>
class continuator_base {
public:
    void init() { ... }

    template <typename T>
    void emit(T&& value) const
    {
        std::invoke(m_continuation, FWD(value));
    }

    void notify_ended() const { ... }

protected:
    Cont m_continuation;
};
```

Invoke

```
std::invoke(function, arg1, arg2, ...)
```

For most cases (functions, function objects, lambdas)
equivalent to:

```
function(arg1, arg2, ...)
```

But it can also invoke class member functions:

```
arg1.function(arg2, ...)
```

Source

```
template <typename Cont>
class values_node: public continuator_base<Cont>, non_copyable {
    ...

    void init()
    {
        base::init();


        for (auto&& value: m_values) {
            base::emit(value);
        }

        m_values.clear();


        base::notify_ended();
    }

    ...
};
```

Source




Creating a source




As far as the rest of the system is concerned, this node **creates** the messages.

Creating a sink



Creating a transformation



Creating a transformation

```
template <typename Cont>
class transform_node: public continuator_base<Cont>, non_copyable {
public:
    ...

    template <typename T>
    void operator() (T&& value) const
    {
        base::emit(std::invoke(m_transformation, FWD(value)));
    }

private:
    Traf m_transformation;
};
```

Filtering

```
template <typename Cont>
class filter_node: public continuator_base<Cont>, non_copyable {
public:
    ...

    template <typename T>
    void operator() (T&& value) const
    {
        if (std::invoke(m_predicate, value) {
            base::emit(FWD(value));
        }
    }

private:
    Predicate m_predicate;
};
```

PIPES

Connecting the components

```
auto pipeline = source | transform(...)
                  | filter(...)
                  | join
                  | transform(...)
                  ...
```

AST

When the pipeline AST is completed and evaluated, then we get proper pipeline nodes.

Associativity

```
auto my_transformation = filter(...)
                        | join;

auto pipeline = source | transform(...)
                        | my_transformation
                        | transform(...)
                        ...
```

Different pipes

- Composing two transformations
`transform(...) | filter(...)`
- Connecting a source to the continuation
`source | transform(...)`
- Ending the pipeline
`filter(...) | sink`
- Closing the pipeline
`source | sink`

Different pipes

```
struct source_node_tag {};  
struct sink_node_tag {};  
struct transformation_node_tag {};  
  
template <typename T>  
class values {  
public:  
    using node_category = source_node_tag;  
  
};
```

Different pipes

```
struct source_node_tag {};  
struct sink_node_tag {};  
struct transformation_node_tag {};  
  
template <typename T>  
class filter {  
public:  
    using node_category = transformation_node_tag;  
  
};
```


Concepts

```
template <typename Left, typename Right>
auto operator| (Left&& left, Right&& right)
{
    return ...;
}
```

Problem: Defines operator| on all types

Concepts

```
template < typename Node
           , typename Category = detected_t<node_category, Node>
           >
f_concept is_node( )
{
    if constexpr ( !is_detected_v<node_category, Node> ) {
        return false;
    } else if constexpr ( std:::is_same_v<void, Category> ) {
        return false;
    } else {
        return true;
    }
}
```

Note: Proper concepts will come in C++20. This uses the detection idiom to simulate concepts.

Concepts

```
template < typename Left
           , typename Right
           , f_require(
               is_node<Left> && is_node<Right>
           )
           >
auto operator| (Left&& left, Right&& right)
{
    return ...;
}
```

Concepts

```
if constexpr (!is_source<Left> && !is_sink<Right>) {  
    ...  
} else if constexpr (is_source<Left> && !is_sink<Right>) {  
    ...  
} else if constexpr (!is_source<Left> && is_sink<Right>) {  
    ...  
} else {  
    ...  
}
```

Concepts

```
if constexpr (!is_source<Left> && !is_sink<Right>) {  
    ...  
} else if constexpr (is_source<Left> && !is_sink<Right>) {  
    ... transformed source is also a source  
} else if constexpr (!is_source<Left> && is_sink<Right>) {  
    ...  
} else {  
    ...  
}
```


Concepts

```
if constexpr (!is_source<Left> && !is_sink<Right>) {  
    ...  
} else if constexpr (is_source<Left> && !is_sink<Right>) {  
    ...  
} else if constexpr (!is_source<Left> && is_sink<Right>) {  
    ... transformation with a sink is a sink  
} else {  
    ...  
}
```


Concepts

```
if constexpr (!is_source<Left> && !is_sink<Right>) {  
    ...  
} else if constexpr (is_source<Left> && !is_sink<Right>) {  
    ...  
} else if constexpr (!is_source<Left> && is_sink<Right>) {  
    ...  
} else {  
    ... the pipeline is complete, evaluate it  
}
```

Pipes


```
auto new_source = source | transform(...) | filter(...);  
auto new_trafo = transform(...) | filter(...);  
auto new_sink = filter(...) | sink;  
auto pipeline = new_source | new_trafo | new_sink;
```

Pipes


- Easy implementation of reactive systems
- Transparent handling of synchronous vs asynchronous transformations
- High-enough abstraction for more powerful idioms

GOING POSTAL


Functional thinking – data transformation



Functional thinking – data transformation



Functional thinking – data transformation



Distributed pipelines

```
auto pipeline =
    system_cmd("ping"s, "localhost"s)

    | transform(string_to_upper)

    // Parse the ping output
    | transform([] (std::string&& value) {
        const auto pos = value.find_last_of('=');
        return std::make_pair(std::move(value), pos);
    })

    // Extract the ping time from the output
    | transform([] (std::pair<std::string, size_t>&& pair) {
        auto [ value, pos ] = pair;
        return pos == std::string::npos
            ? std::move(value)
            : std::string(value.cbegin() + pos + 1, value.cend());
    })

    // Remove slow pings
    | filter([] (const std::string& value) {
        return value < "0.145"s;
    })

    // Print out the ping info
    | sink{cout};
```

Distributed pipelines

```
auto pipeline =
    system_cmd("ping"s, "localhost"s)

    | transform(string_to_upper)

    | voy_bridge(frontend_to_backend_1)

    | transform([] (std::string&& value) {
        const auto pos = value.find_last_of('=');
        return std::make_pair(std::move(value), pos);
    })

    | transform([] (std::pair<std::string, size_t>&& pair) {
        auto [ value, pos ] = pair;
        return pos == std::string::npos
            ? std::move(value)
            : std::string(value.cbegin() + pos + 1, value.cend());
    })

    | voy_bridge(backend_1_to_backend_2)

    | filter([] (const std::string& value) {
        return value < "0.145"s;
    })

    | voy_bridge(backend_2_to_frontend)

    | sink{cout};
```

Iterators
○○○○○

Ranges
○○○○○○○○○○○○○○○○

Push
○○○○○○○○○○○○○○○○

Pipes
○○○○○○○○○○○○○○○○

Going postal
○○○○●○


Live demo

Answers? Questions! Questions? Answers!

Kudos (in chronological order):

Friends at **KDE**

Saša Malkov and **Zoltán Porkoláb**
спасибо **Сергею и двум Антонам**



MEAP – Manning Early Access Program
Functional Programming in C++
cukic.co/to/fp-in-cpp

