

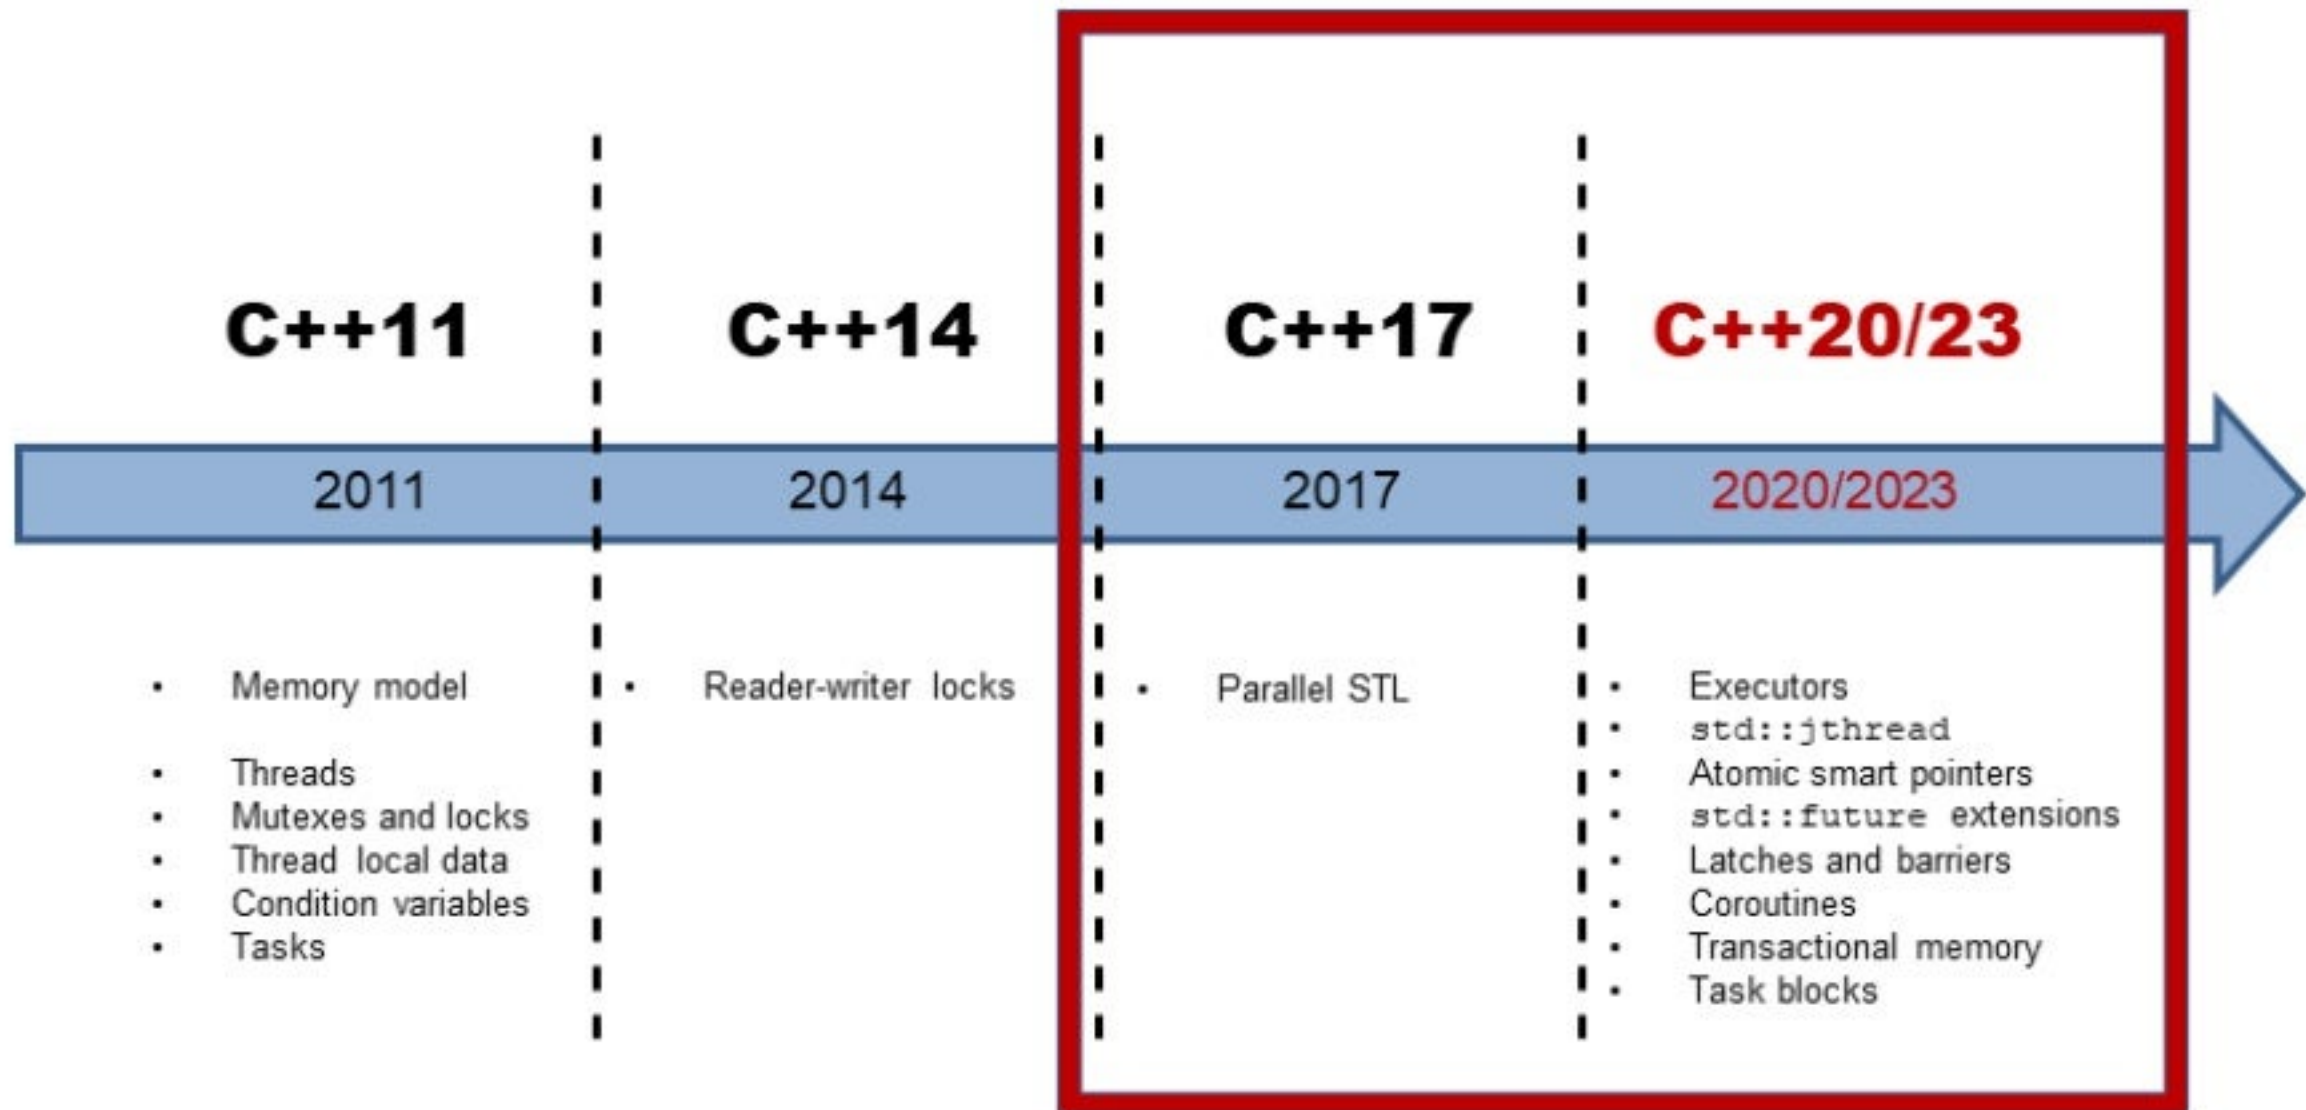
# Concurrency and Parallelism with C++17 and C++20/23

Rainer Grimm

Training, Coaching and  
Technology Consulting

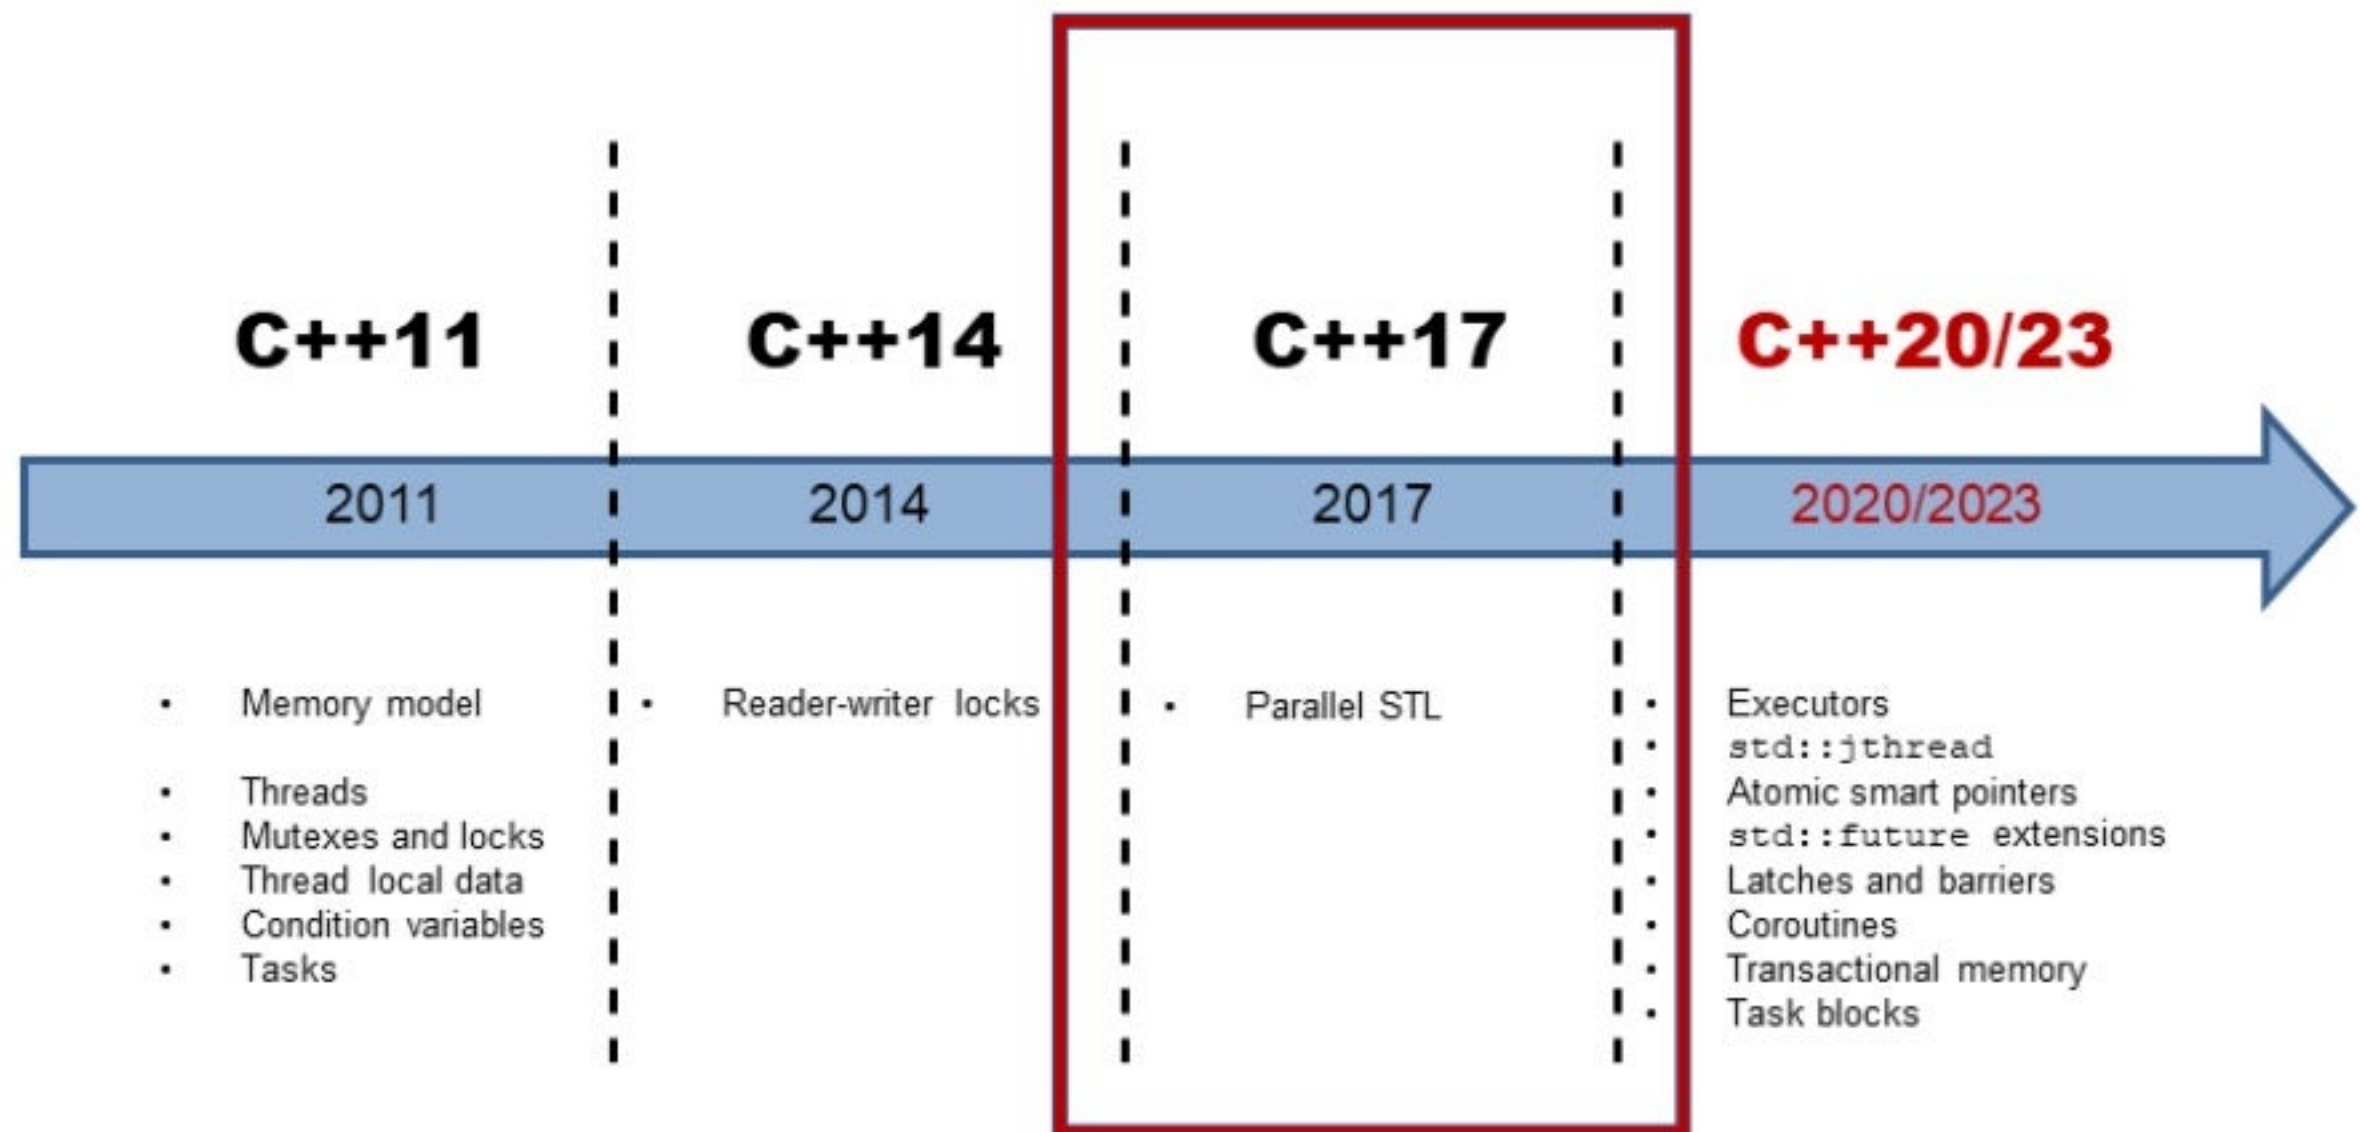
[www.ModernesCpp.de](http://www.ModernesCpp.de)

# Concurrency and Parallelism in C++





# Concurrency and Parallelism in C++17



# Parallel STL

You can choose the execution policy of an algorithm.

- Execution policies

- `std::execution::seq`

- Sequential in one thread

- `std::execution::par`

- Parallel

- `std::execution::par_unseq`

- Parallel and vectorised ➡ SIMD

# Parallel STL

```
const int SIZE = 8;
int vec[]={1, 2 , 3, 4, 5, 6, 7, 8};
int res[SIZE] = {0,};

int main(){
    for (int i= 0; i < SIZE; ++i){
        res[i] = vec[i] + 5;
    }
}
```

## Not vectorised

```
movslq -8(%rbp), %rax
movl   vec(,%rax,4), %ecx
addl   $5, %ecx
movslq -8(%rbp), %rax
movl   %ecx, res(,%rax,4)
```

## Vectorised

```
movdqa .LCPI0_0(%rip), %xmm0 # xmm0 = [5,5,5,5]
movdqa vec(%rip), %xmm1
padd   %xmm0, %xmm1
movdqa %xmm1, res(%rip)
padd   vec+16(%rip), %xmm0
movdqa %xmm0, res+16(%rip)
xorl   %eax, %eax
```

# Parallel STL

```
using namespace std;
vector<int> vec = {1, 2, 3, 4, 5, .... }

sort(vec.begin(), vec.end());           // sequential as ever

sort(execution::seq, vec.begin(), vec.end());           // sequential
sort(execution::par, vec.begin(), vec.end());           // parallel
sort(execution::par_unseq, vec.begin(), vec.end());     // par + vec
```



# Parallel STL

adjacent\_difference, adjacent\_find, all\_of, any\_of, copy, copy\_if, copy\_n, count, count\_if, equal, **exclusive\_scan**, fill, fill\_n, find, find\_end, find\_first\_of, find\_if, find\_if\_not, **for\_each**, **for\_each\_n**, generate, generate\_n, includes, **inclusive\_scan**, inner\_product, inplace\_merge, is\_heap, is\_heap\_until, is\_partitioned, is\_sorted, is\_sorted\_until, lexicographical\_compare, max\_element, merge, min\_element, minmax\_element, mismatch, move, none\_of, nth\_element, partial\_sort, partial\_sort\_copy, partition, partition\_copy, **reduce**, remove, remove\_copy, remove\_copy\_if, remove\_if, replace, replace\_copy, replace\_copy\_if, replace\_if, reverse, reverse\_copy, rotate, rotate\_copy, search, search\_n, set\_difference, set\_intersection, set\_symmetric\_difference, set\_union, sort, stable\_partition, stable\_sort, swap\_ranges, transform, **transform\_exclusive\_scan**, **transform\_inclusive\_scan**, **transform\_reduce**, uninitialized\_copy, uninitialized\_copy\_n, uninitialized\_fill, uninitialized\_fill\_n, unique, unique\_copy

# Parallel STL

`std::transform_reduce`

- Haskell's function `map` is called `std::transform` in C++
- `std::transform_reduce`  `std::map_reduce`

```
std::vector<std::string> strVec{"Only", "for", "testing", "purpose"};

std::size_t res = std::transform_reduce(std::execution::par,
                                       strVec.begin(), strVec.end(), 0,
                                       [] (std::size_t a, std::size_t b) { return a + b; },
                                       [] (std::string s) { return s.length(); });

std::cout << res;    // 21
```



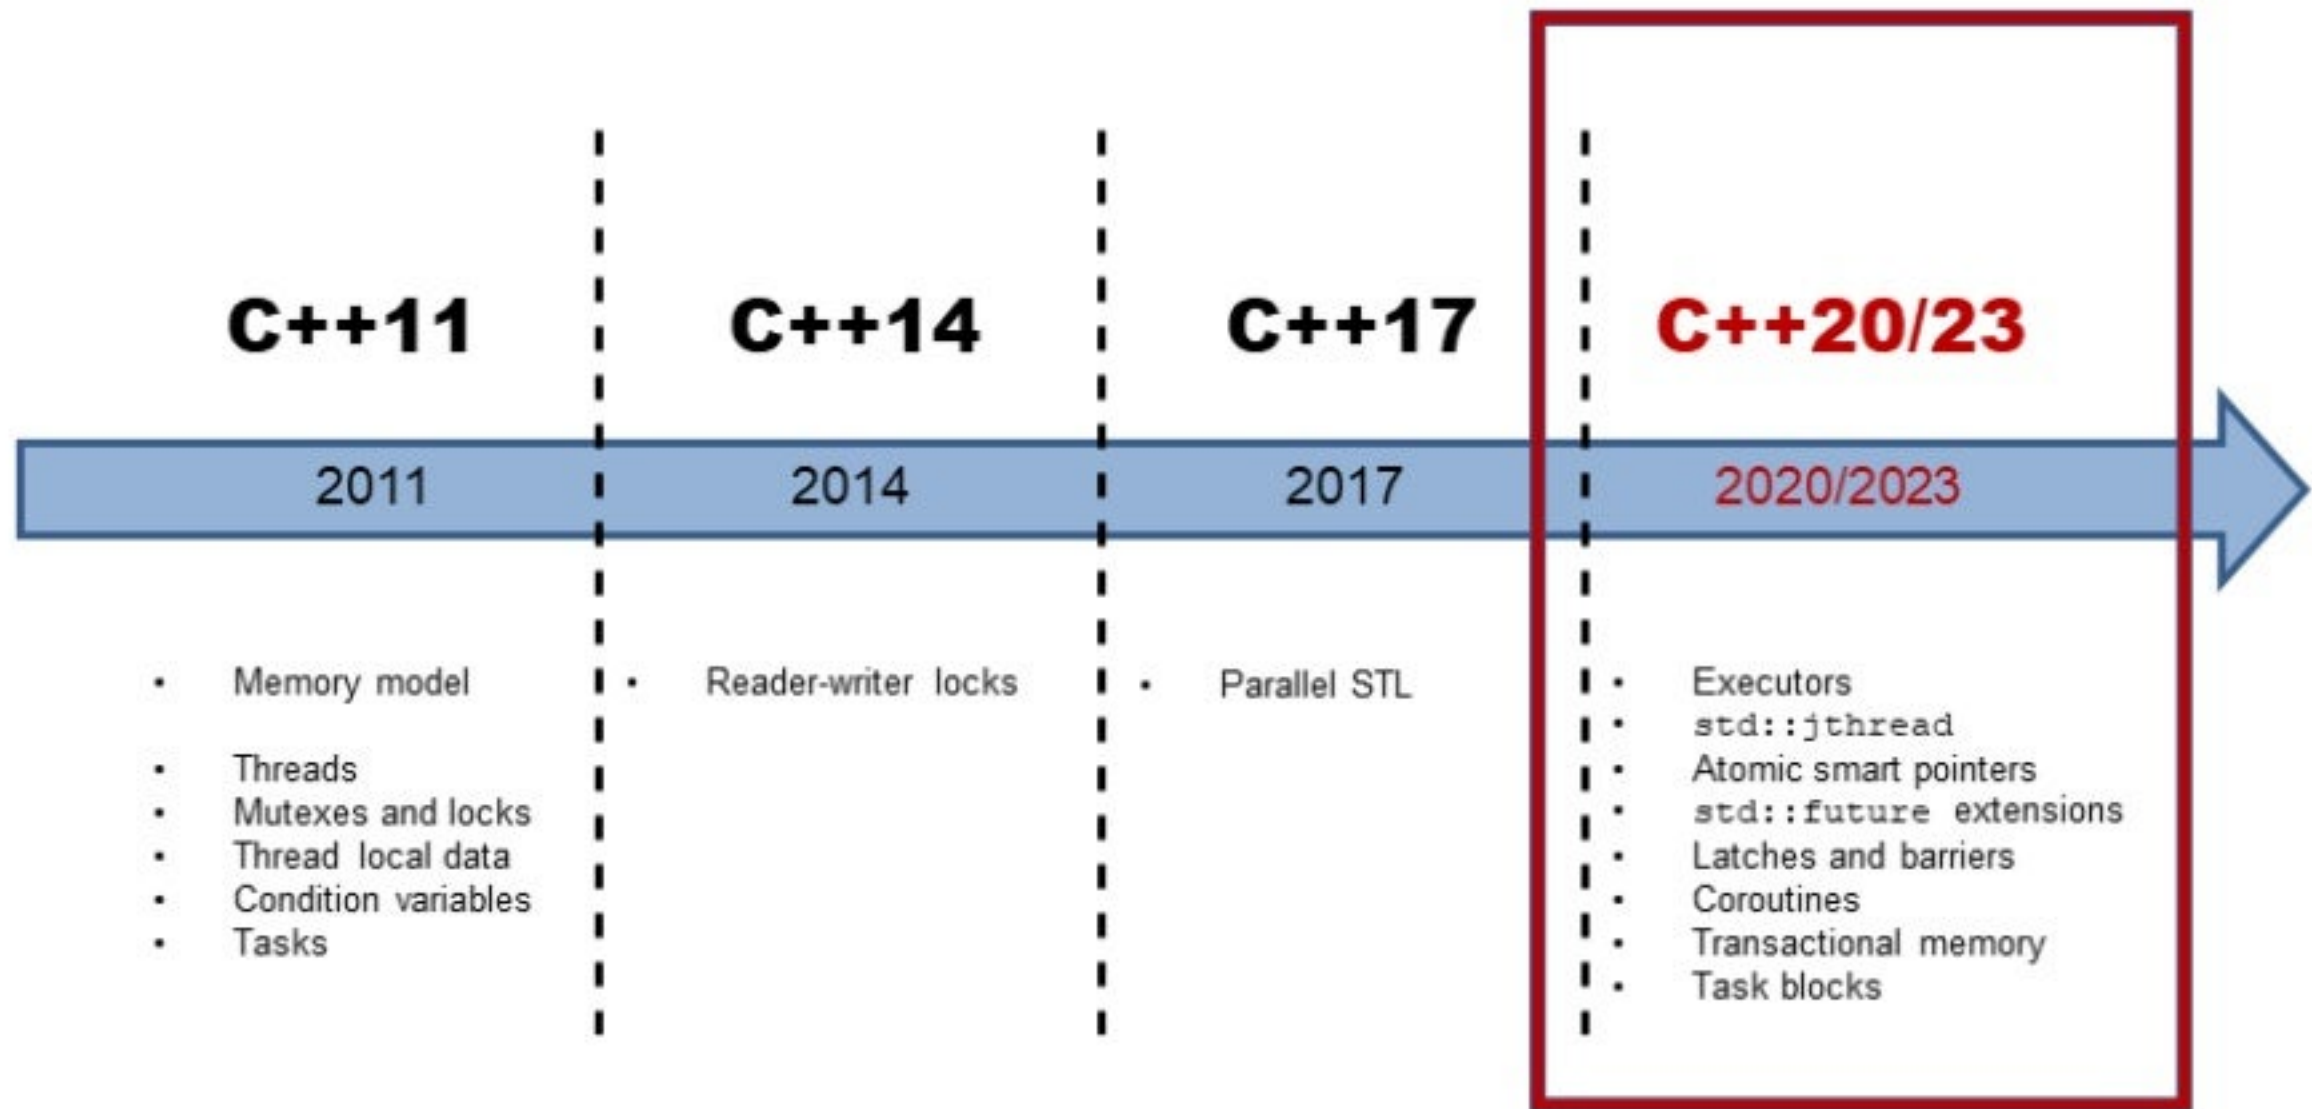
# Parallel STL

- Danger of data races or deadlocks

```
int numComp = 0;
vector<int> vec = {1, 3, 8, 9, 10};
sort(execution::par, vec.begin(), vec.end(),
      [&numComp](int fir, int sec){ numComp++; return fir < sec; }
);
```

➡ The access to **numComp** has to be atomic.

# Concurrency and Parallelism in C++20/23



# Executors

Executors are the basic building block for execution in C++.

- They fulfil a similar role for execution such as allocators for allocation.

An executor consists of a set of rules for a callables:

- **Where:** run on a internal or external processor
- **When:** run immediately or later
- **How:** run on a CPU or GPU



# Executors

- Using an executor

```
my_executor_type my_executor = ... ;  
auto future = std::async(my_executor, []{  
    std::cout << "Hello world " << std::endl; }  
);
```

```
std::for_each(std::execution::par.on(my_executor),  
    data.begin(), data.end(), func);
```

- Obtaining an executor

```
static_thread_pool pool(4);  
auto exec = pool.executor();  
task1 = long_running_task(exec);
```

# Executors

An executor provides one or more execution functions for creating a callable.

Name	Cardinality	Direction
execute	single	oneway
twoway_execute	single	twoway
then_execute	single	then
bulk_execute	bulk	oneway
bulk_twoway_execute	bulk	twoway
bulk_then_execute	bulk	then

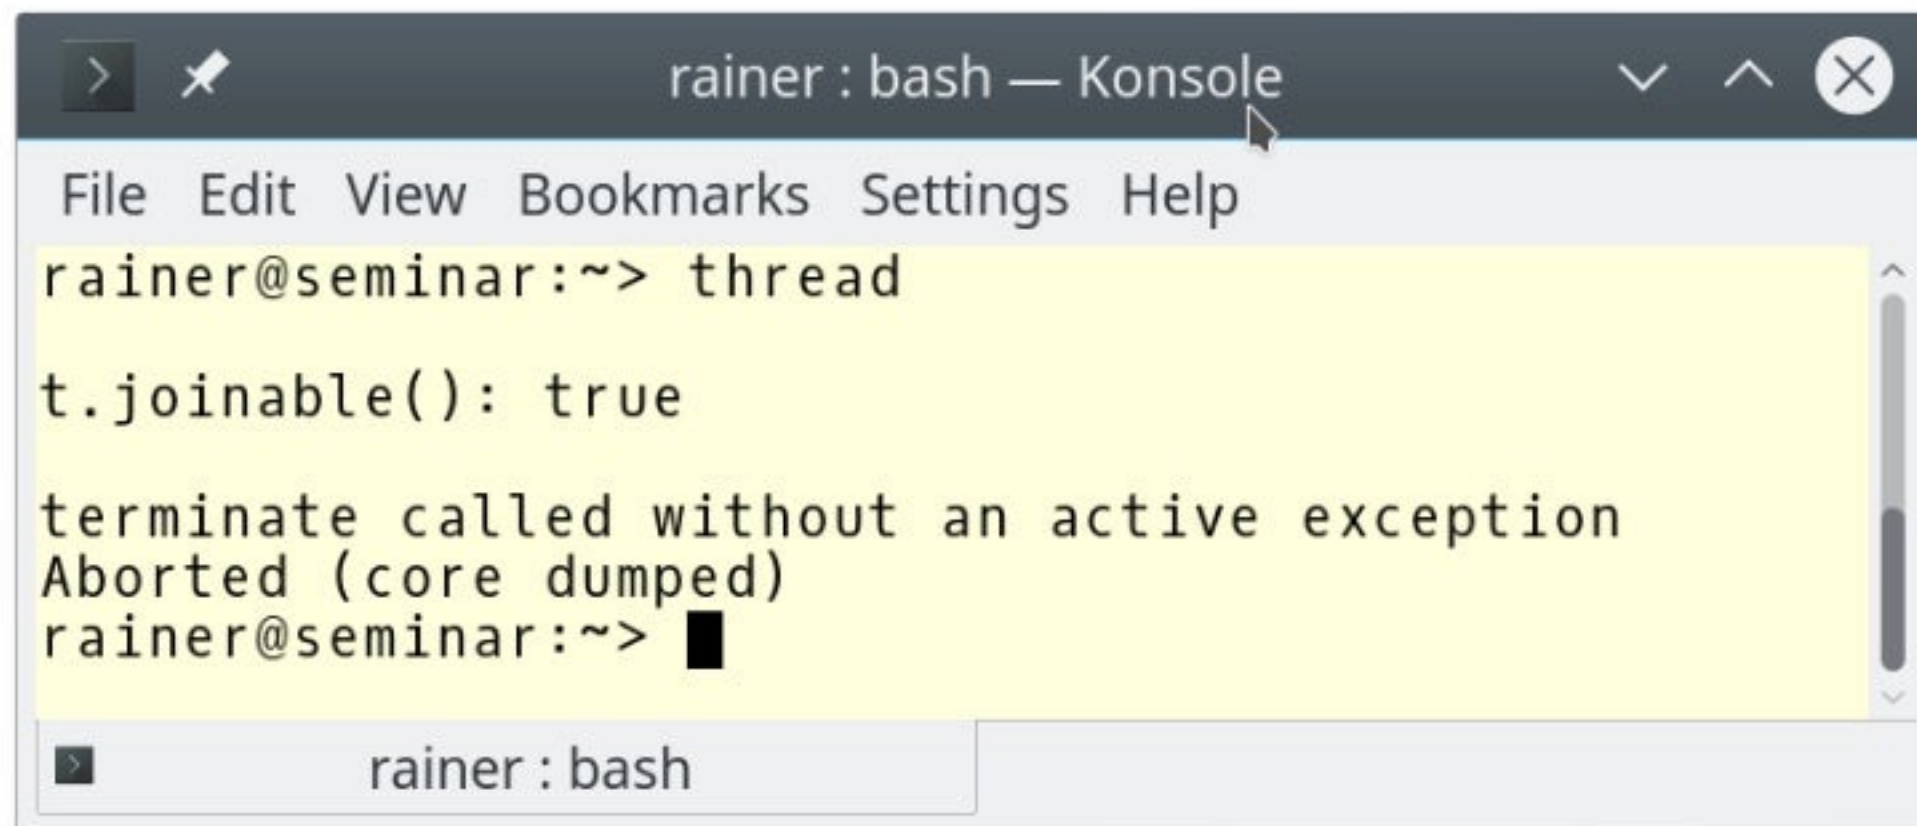
**Cardinality:** Creation of one execution agent or a group of execution agents.

**Direction:** Directions of the execution.

# std::jthread

**Problem:** `std::thread` throws `std::terminate` in its destructor if still joinable.

```
std::thread t{[] { std::cout << "New thread"; }};  
std::cout << "t.joinable(): " << t.joinable();
```



The screenshot shows a terminal window titled "rainer : bash — Konsole". The terminal output is as follows:

```
File Edit View Bookmarks Settings Help  
rainer@seminar:~> thread  
t.joinable(): true  
terminate called without an active exception  
Aborted (core dumped)  
rainer@seminar:~> █
```

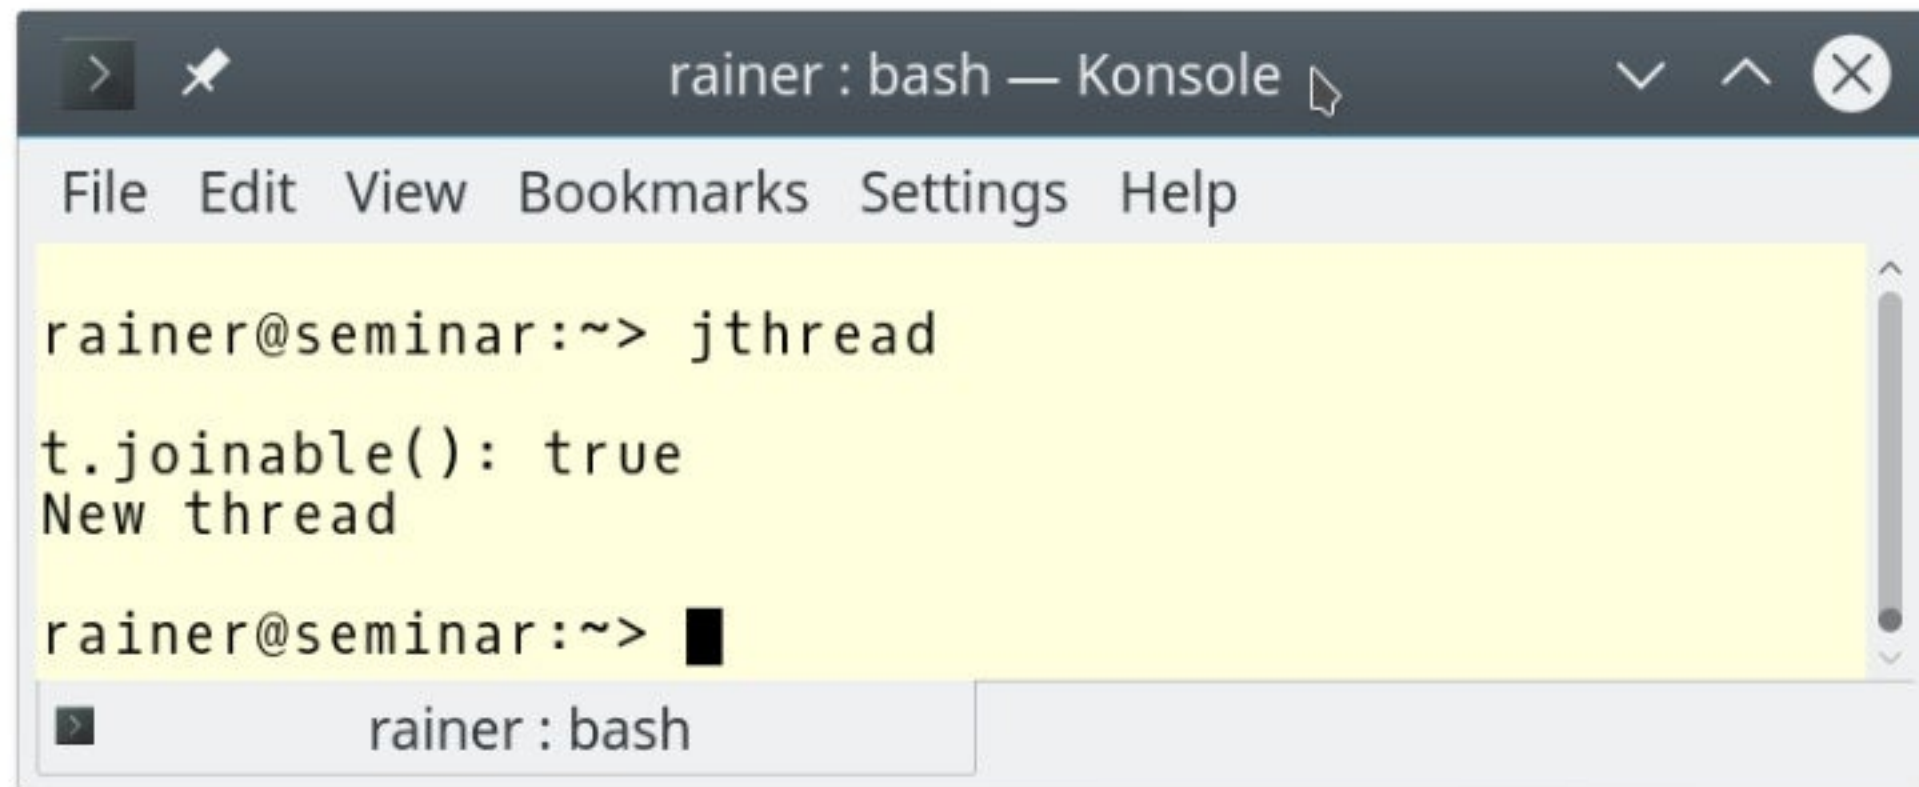
The terminal window has a menu bar with "File", "Edit", "View", "Bookmarks", "Settings", and "Help". The status bar at the bottom shows "rainer : bash".



`std::jthread`

**Solution:** `std::jthread` joins automatically at the end of its scope.

```
std::jthread t{[] { std::cout << "New thread"; }};  
std::cout << "t.joinable(): " << t.joinable();
```



The screenshot shows a terminal window titled "rainer : bash — Konsole". The terminal output is as follows:

```
rainer@seminar:~> jthread  
t.joinable(): true  
New thread  
rainer@seminar:~> █
```

The output demonstrates that the thread is joinable and prints "New thread" before the program ends, indicating automatic joining.

## `std::jthread`

- Instances of `std::jthread` can be interrupted

### Receiver

- Explicit check:
  - **`is_interrupted`**: yields, when an interrupt was signalled
  - `std::condition_variable` **`wait`** variations with predicate

### Sender

- **`interrupt`**: signals an interrupt (and returns whether an interrupt was signaled before)

# std::jthread

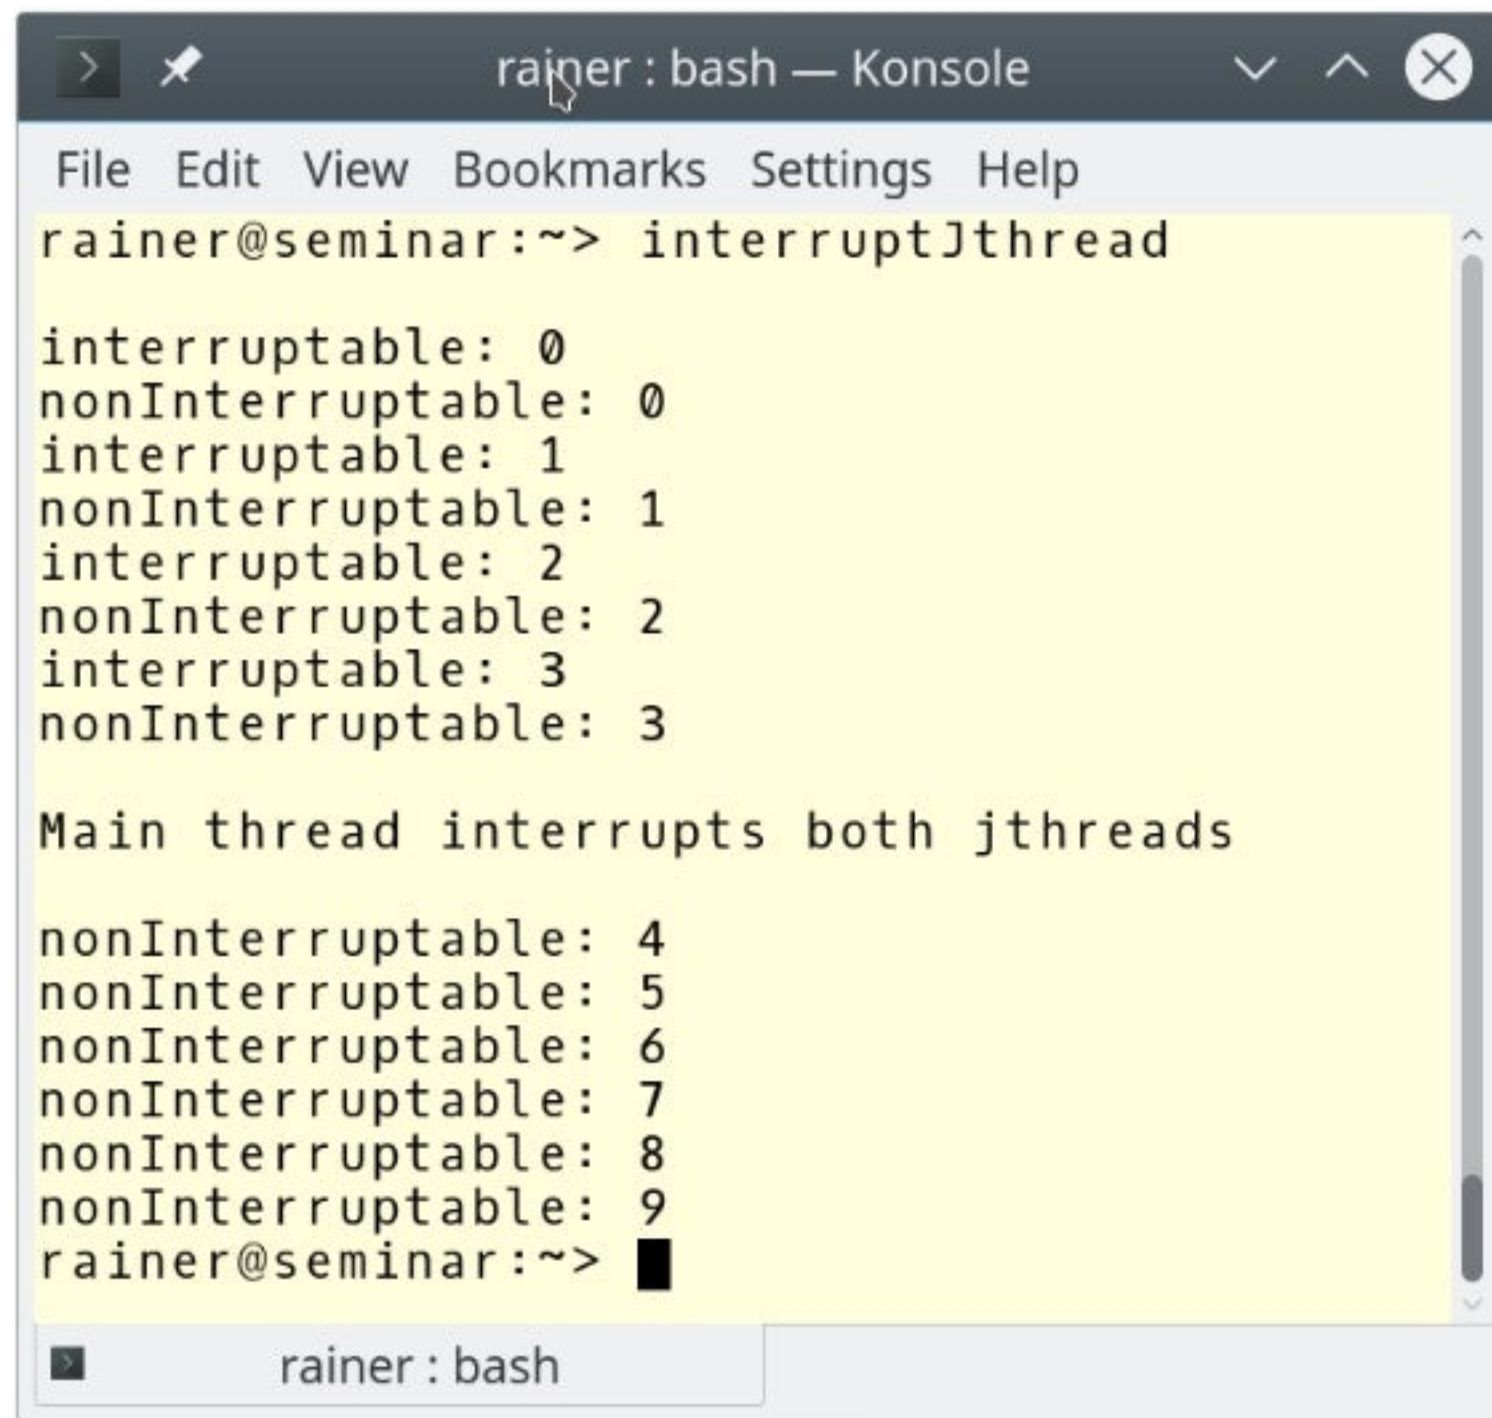
```
jthread nonInterruptable([]{  
    int counter{0};  
    while (counter < 10){  
        this_thread::sleep_for(0.2s);  
        cerr << "nonInterruptable: "  
            << counter << endl;  
        ++counter;  
    }  
});
```

```
jthread interruptable([](interrupt_token  
itoken){  
    int counter{0};  
    while (counter < 10){  
        this_thread::sleep_for(0.2s);  
        if (itoken.is_interrupted()) return;  
        cerr << "interruptable: "  
            << counter << endl;  
        ++counter;  
    }  
});
```

```
this_thread::sleep_for(1s);  
cerr << endl;  
cerr << "Main thread interrupts both jthreads" << endl;  
nonInterruptable.interrupt();  
interruptable.interrupt();  
cout << endl;
```



# std::jthread



```
rainer : bash — Konsole
File Edit View Bookmarks Settings Help
rainer@seminar:~> interruptJthread

interruptable: 0
nonInterruptable: 0
interruptable: 1
nonInterruptable: 1
interruptable: 2
nonInterruptable: 2
interruptable: 3
nonInterruptable: 3

Main thread interrupts both jthreads

nonInterruptable: 4
nonInterruptable: 5
nonInterruptable: 6
nonInterruptable: 7
nonInterruptable: 8
nonInterruptable: 9
rainer@seminar:~>
```



# Atomic Smart Pointers

C++11 has a `std::shared_ptr` for shared ownership.

- General Rule:
  - You should use smart pointers.
- But:
  - The managing of the control block and the deletion of the resource is thread-safe. The access to the resource is not thread-safe.
- ➡ **Tony Van Eerd: Forget what you learned in Kindergarten. Stop sharing.**
- *Solution:*
  - `std::atomic_shared_ptr`
  - `std::atomic_weak_ptr`

# Atomic Smart Pointer

## 3 Reasons

- Consistency
  - `std::shared_ptr` is the only non-atomic data type for which atomic operations exists.
- Correctness
  - The correct usage of atomic operations is just based on the discipline of the user.  extremely error-prone
  - `std::atomic_store(&sharPtr, localPtr)  sharPtr = localPtr`
- Performance
  - `std::shared_ptr` has to be design for the special use-case.



# Atomic Smart Pointer

```
template<typename T> class concurrent_stack {
    struct Node { T t; shared_ptr<Node> next; };
    atomic_shared_ptr<Node> head;
    // in C++11: remove "atomic_" and remember to use the special
    // functions every time you touch the variable
    concurrent_stack( concurrent_stack &) =delete;
    void operator=(concurrent_stack&) =delete;

public:
    concurrent_stack() =default;
    ~concurrent_stack() =default;
    class reference {
        shared_ptr<Node> p;
    public:
        reference(shared_ptr<Node> p_) : p{p_} { }
        T& operator* () { return p->t; }
        T* operator->() { return &p->t; }
    };

    auto find( T t ) const {
        auto p = head.load(); // in C++11: atomic_load(&head)
        while( p && p->t != t )
            p = p->next;
        return reference(move(p));
    }
    auto front() const {
        return reference(head); // in C++11: atomic_load(&head)
    }
    void push_front( T t ) {
        auto p = make_shared<Node>();
        p->t = t;
        p->next = head; // in C++11: atomic_load(&head)
        while( !head.compare_exchange_weak(p->next, p) ){ }
        // in C++11: atomic_compare_exchange_weak(&head, &p->next, p);
    }
    void pop_front() {
        auto p = head.load();
        while( p && !head.compare_exchange_weak(p, p->next) ){ }
        // in C++11: atomic_compare_exchange_weak(&head, &p, p->next);
    }
};
```

# Atomic Smart Pointers

Atomic smart pointers are part of ISO C++ standard.

- `std::atomic_shared_ptr`  
➡ `std::atomic<std::shared_ptr<T>>`
- `std::atomic_weak_ptr`  
➡ `std::atomic<std::weak_ptr<T>>`

# std::future Extensions

std::future doesn't support composition

- std::future Improvement ➡ Continuation
  - then: execute the next future if the previous one is done

```
future<int> f1 = async([](){ return 123; });
future<string> f2 = f1.then([](future<int> f){
    return to_string(f.get()); // non-blocking
});
auto myResult = f2.get(); // blocking
```



# std::future Extensions

- **when\_all**: execute the future if all futures are done

```
future<int> futures[] = { async([]() { return intResult(125); }),  
                        async([]() { return intResult(456); }) };  
future<vector<future<int>>> all_f = when_all(begin(futures), end(futures));  
  
vector<future<int>> myResult = all_f.get();  
  
for (auto fut: myResult): fut.get();
```

- **when\_any**: execute the future if one of the futures is done

```
future<int> futures[] = { async([]() { return intResult(125); }),  
                        async([]() { return intResult(456); }) };  
when_any_result<vector<future<int>>> any_f = when_any(begin(futures),  
                                                    end(futures));  
  
future<int> myResult = any_f.futures[any_f.index];  
  
auto myResult = myResult.get();
```

# std::future Extensions

- Disadvantages of the extended futures

- Futures and promises are coupled to `std::thread`.
- Where are `.then` continuations are invoked?
- Passing futures to `.then` continuation is too verbose.

```
std::future f1 = std::async([]{ return 123; });  
std::future f2 = f1.then([](std::future f){  
    return std::to_string(f.get());  
});
```

```
std::future f2 = f1.then(std::to_string);
```

- Future blocks in its destructor.
- Futures und values should be easily composable.

```
bool f(std::string, double, int);  
std::future<std::string> a = /* ... */;  
std::future<int> c = /* ... */;  
std::future<bool> d2 = when_all(a, 3.14, c).then(f);  
// f(a.get(), 3.14, c.get())
```

# Latches and Barriers

C++ has no semaphor → latches and barriers

- Key idea

A thread is waiting at the synchronisation point until the counter becomes zero.

- `latch` is for the one-time use-case

- `count_down_and_wait`: decrements the counter until it becomes zero
- `count_down(n = 0)`: decrements the counter by `n`
- `is_ready`: checks the counter
- `wait`: waits until the counter becomes zero



# Latches and Barriers

- `barrier` can be reused
  - `arrive_and_wait`: waits at the synchronisation point
  - `arrive_and_drop`: removes itself from the synchronisation mechanism
- `flex_barrier` is a reusable and adaptable barrier
  - The constructor gets a callable.
  - The callable will be called in the completion phase.
  - The callable returns a number which stands for the counter in the next iteration.
  - Can change the value of the counter for each iteration.

# Latches and Barriers

```
void doWork(threadpool* pool) {  
    latch completion_latch(NUMBER_TASKS);  
    for (int i = 0; i < NUMBER_TASKS; ++i) {  
        pool->add_task([&] {  
            // perform the work  
            ...  
            completion_latch.count_down();  
        });  
    }  
    // block until all tasks are done  
    completion_latch.wait();  
}
```

# Coroutines

Coroutines are generalised functions that can be suspended and resumed while keeping their state.

- Typical use-case
  - Cooperative Tasks
  - Event loops
  - Infinite data streams
  - Pipelines



# Coroutines

	Function	Coroutine
invoke	<code>func(args)</code>	<code>func(args)</code>
return	<code>return statement</code>	<code>co_return someValue</code>
suspend		<code>co_await someAwaitable</code> <code>co_yield someValue</code>
resume		<code>coroutine_handle&lt;&gt;::resume()</code>

A function is a coroutine if it has a `co_return`, `co_await`, `co_yield` call or if it has a range-based for loop with a `co_await` call.

# Coroutines

```
generator<int> genForNumbers(int begin, int inc= 1){  
    for (int i = begin;; i += inc){  
        co_yield i;  
    }  
}
```

```
int main(){  
    auto numbers = genForNumbers(-10);  
    for (int i = 1; i <= 20; ++i) std::cout << numbers << " ";  
    for (auto n: genForNumbers(0,5)) std::cout << n << " ";  
}
```

➡ **-10 -9 -8 -7 -6 -5 -4 -3 -2 -1 0 1 2 3 4 5 6 7 8 9 10**  
**0 5 10 15 20 25 30 35 40 45 50 55 60 65 70 75 80 85 ....**



# Coroutines

## Blocking

```
Acceptor accept{443};

while (true){
    Socket so = accept.accept(); // block
    auto req = so.read();         // block
    auto resp = handleRequest(req);
    so.write(resp);               // block
}
```

## Waiting

```
Acceptor accept{443};

while (true){
    Socket so = co_await accept.accept();
    auto req = co_await so.read();
    auto resp = handleRequest(req);
    co_await so.write(resp);
}
```



# Transactional Memory

*Transactional Memory* is the idea of transactions from the data base theory applied to software.

- A transaction has the ACID properties without ***Durability***


```
atomic{  
    statement1;  
    statement2;  
    statement3;  
}
```

- **Atomicity:** all or no statement will be performed
- **Consistency:** the system is always in a consistent state
- **Isolation:** a transaction runs total isolation
- **Durability:** the result of a transaction will be stored

# Transactional Memory

- Transactions
  - build a total order
  - feel like a global lock
    - ➔ Optimistic approach  $\neq$  lock
- Workflow

## *Retry*



A transaction stores its initial state.

The transaction will be performed without synchronisation.

The runtime experiences a violation to the initial state.

The transaction will be performed once more.

## *Rollback*



# Transactional Memory

- Two forms
  - `synchronized blocks`
    - *relaxed* transaction
    - are not transaction in the pure sense
    - ➔ can have `transaction-unsafe code`
  - `atomic blocks`
    - atomic blocks
    - are available in three variations
    - ➔ can only execute `transaction-safe code`



# Transactional Memory

```
int i = 0;

void inc() {
    synchronized{
        cout << ++i << " ,";
    }
}

vector<thread> vecSyn(10);
for(auto& t: vecSyn)
    t = thread([]{ for(int n = 0; n < 10; ++n) inc(); });
```



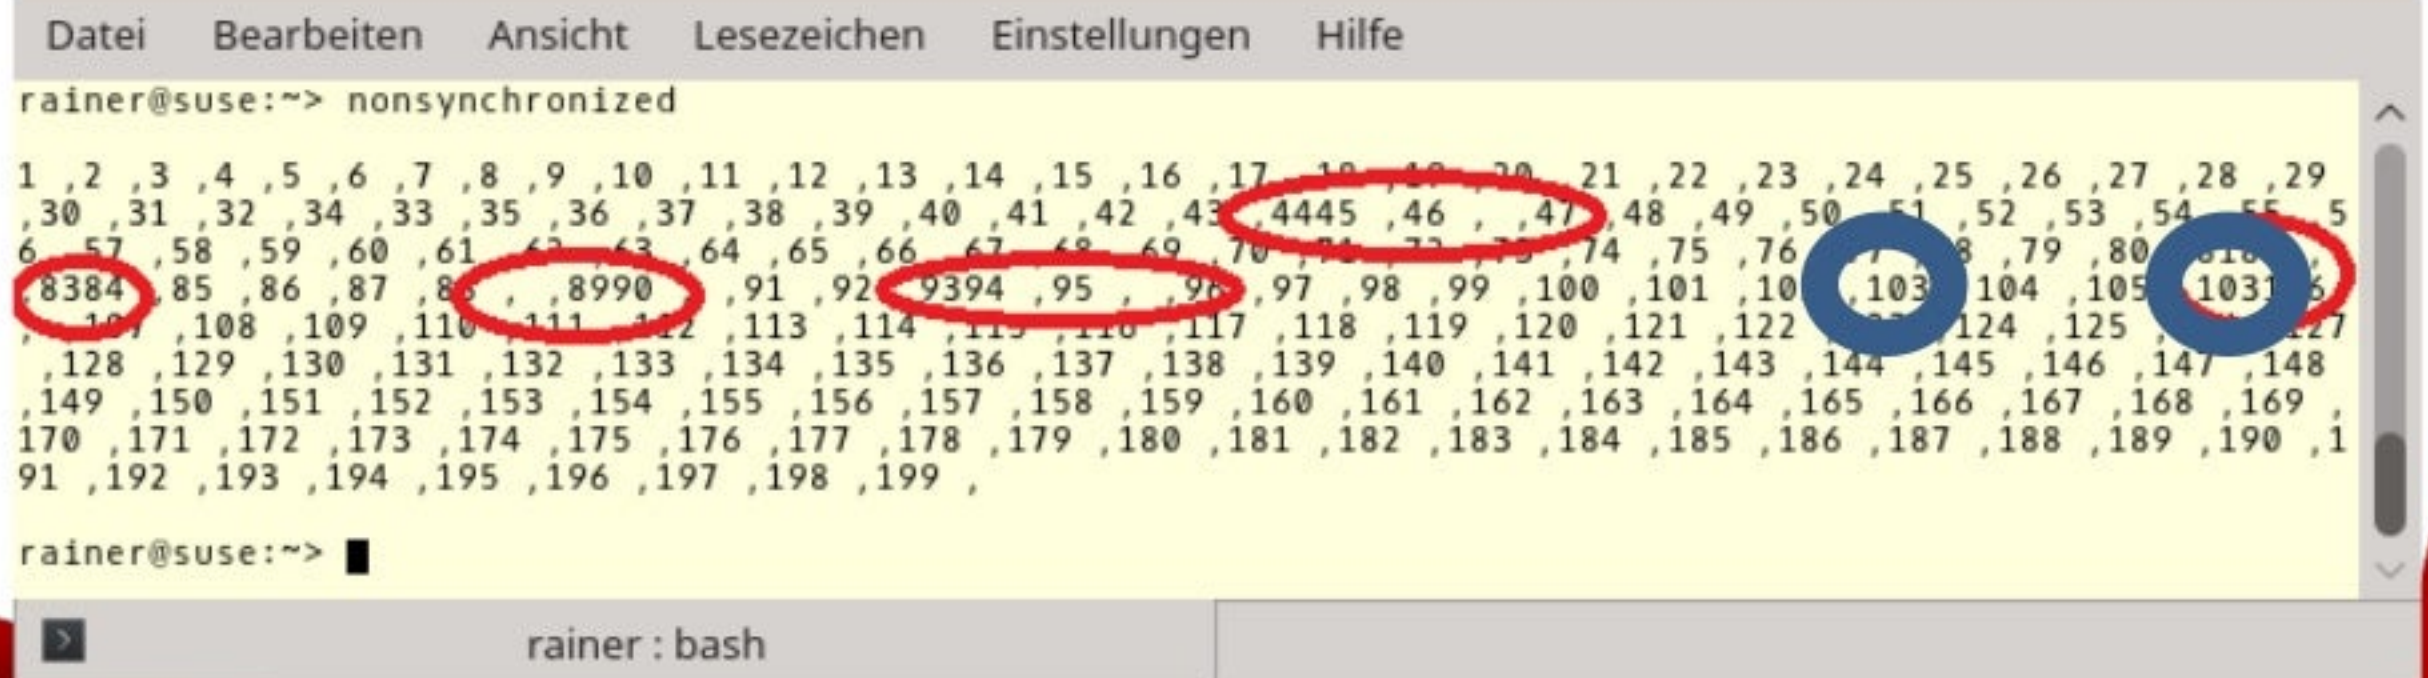
```
Datei Bearbeiten Ansicht Lesezeichen Einstellungen Hilfe
rainer@suse:~> synchronized
1 , 2 , 3 , 4 , 5 , 6 , 7 , 8 , 9 , 10 , 11 , 12 , 13 , 14 , 15 , 16 , 17 , 18 , 19 , 20 , 21 , 22 , 23 , 24 , 25 , 26 , 27 , 28 , 29
0 , 31 , 32 , 33 , 34 , 35 , 36 , 37 , 38 , 39 , 40 , 41 , 42 , 43 , 44 , 45 , 46 , 47 , 48 , 49 , 50 , 51 , 52 , 53 , 54 , 55 , 56
7 , 58 , 59 , 60 , 61 , 62 , 63 , 64 , 65 , 66 , 67 , 68 , 69 , 70 , 71 , 72 , 73 , 74 , 75 , 76 , 77 , 78 , 79 , 80 , 81 , 82 , 83
4 , 85 , 86 , 87 , 88 , 89 , 90 , 91 , 92 , 93 , 94 , 95 , 96 , 97 , 98 , 99 , 100 ,
rainer@suse:~> █
```

rainer : bash

# Transactional Memory

```
void inc() {  
    synchronized{  
        std::cout << ++i << " ,";  
        this_thread::sleep_for(1ns);  
    }  
}
```


```
vector<thread> vecSyn(10), vecUnsyn(10);  
for(auto& t: vecSyn)  
    t= thread[]{ for(int n = 0; n < 10; ++n) inc(); });  
for(auto& t: vecUnsyn)  
    t= thread[]{ for(int n = 0; n < 10; ++n) cout << ++i << " ,"; });
```



```
Datei Bearbeiten Ansicht Lesezeichen Einstellungen Hilfe  
rainer@suse:~> nonsynchronized  
1 , 2 , 3 , 4 , 5 , 6 , 7 , 8 , 9 , 10 , 11 , 12 , 13 , 14 , 15 , 16 , 17 , 18 , 19 , 20 , 21 , 22 , 23 , 24 , 25 , 26 , 27 , 28 , 29 ,  
30 , 31 , 32 , 33 , 34 , 35 , 36 , 37 , 38 , 39 , 40 , 41 , 42 , 43 , 44 , 45 , 46 , 47 , 48 , 49 , 50 , 51 , 52 , 53 , 54 , 55 , 56 ,  
57 , 58 , 59 , 60 , 61 , 62 , 63 , 64 , 65 , 66 , 67 , 68 , 69 , 70 , 71 , 72 , 73 , 74 , 75 , 76 , 77 , 78 , 79 , 80 , 81 , 82 ,  
83 , 84 , 85 , 86 , 87 , 88 , 89 , 90 , 91 , 92 , 93 , 94 , 95 , 96 , 97 , 98 , 99 , 100 , 101 , 102 , 103 , 104 , 105 , 106 , 107 ,  
108 , 109 , 110 , 111 , 112 , 113 , 114 , 115 , 116 , 117 , 118 , 119 , 120 , 121 , 122 , 123 , 124 , 125 , 126 , 127 ,  
128 , 129 , 130 , 131 , 132 , 133 , 134 , 135 , 136 , 137 , 138 , 139 , 140 , 141 , 142 , 143 , 144 , 145 , 146 , 147 , 148 ,  
149 , 150 , 151 , 152 , 153 , 154 , 155 , 156 , 157 , 158 , 159 , 160 , 161 , 162 , 163 , 164 , 165 , 166 , 167 , 168 , 169 ,  
170 , 171 , 172 , 173 , 174 , 175 , 176 , 177 , 178 , 179 , 180 , 181 , 182 , 183 , 184 , 185 , 186 , 187 , 188 , 189 , 190 , 191 ,  
192 , 193 , 194 , 195 , 196 , 197 , 198 , 199 ,  
rainer@suse:~> █  
rainer : bash
```

# Task Blocks

Fork-join parallelism with task blocks.





# Task Blocks

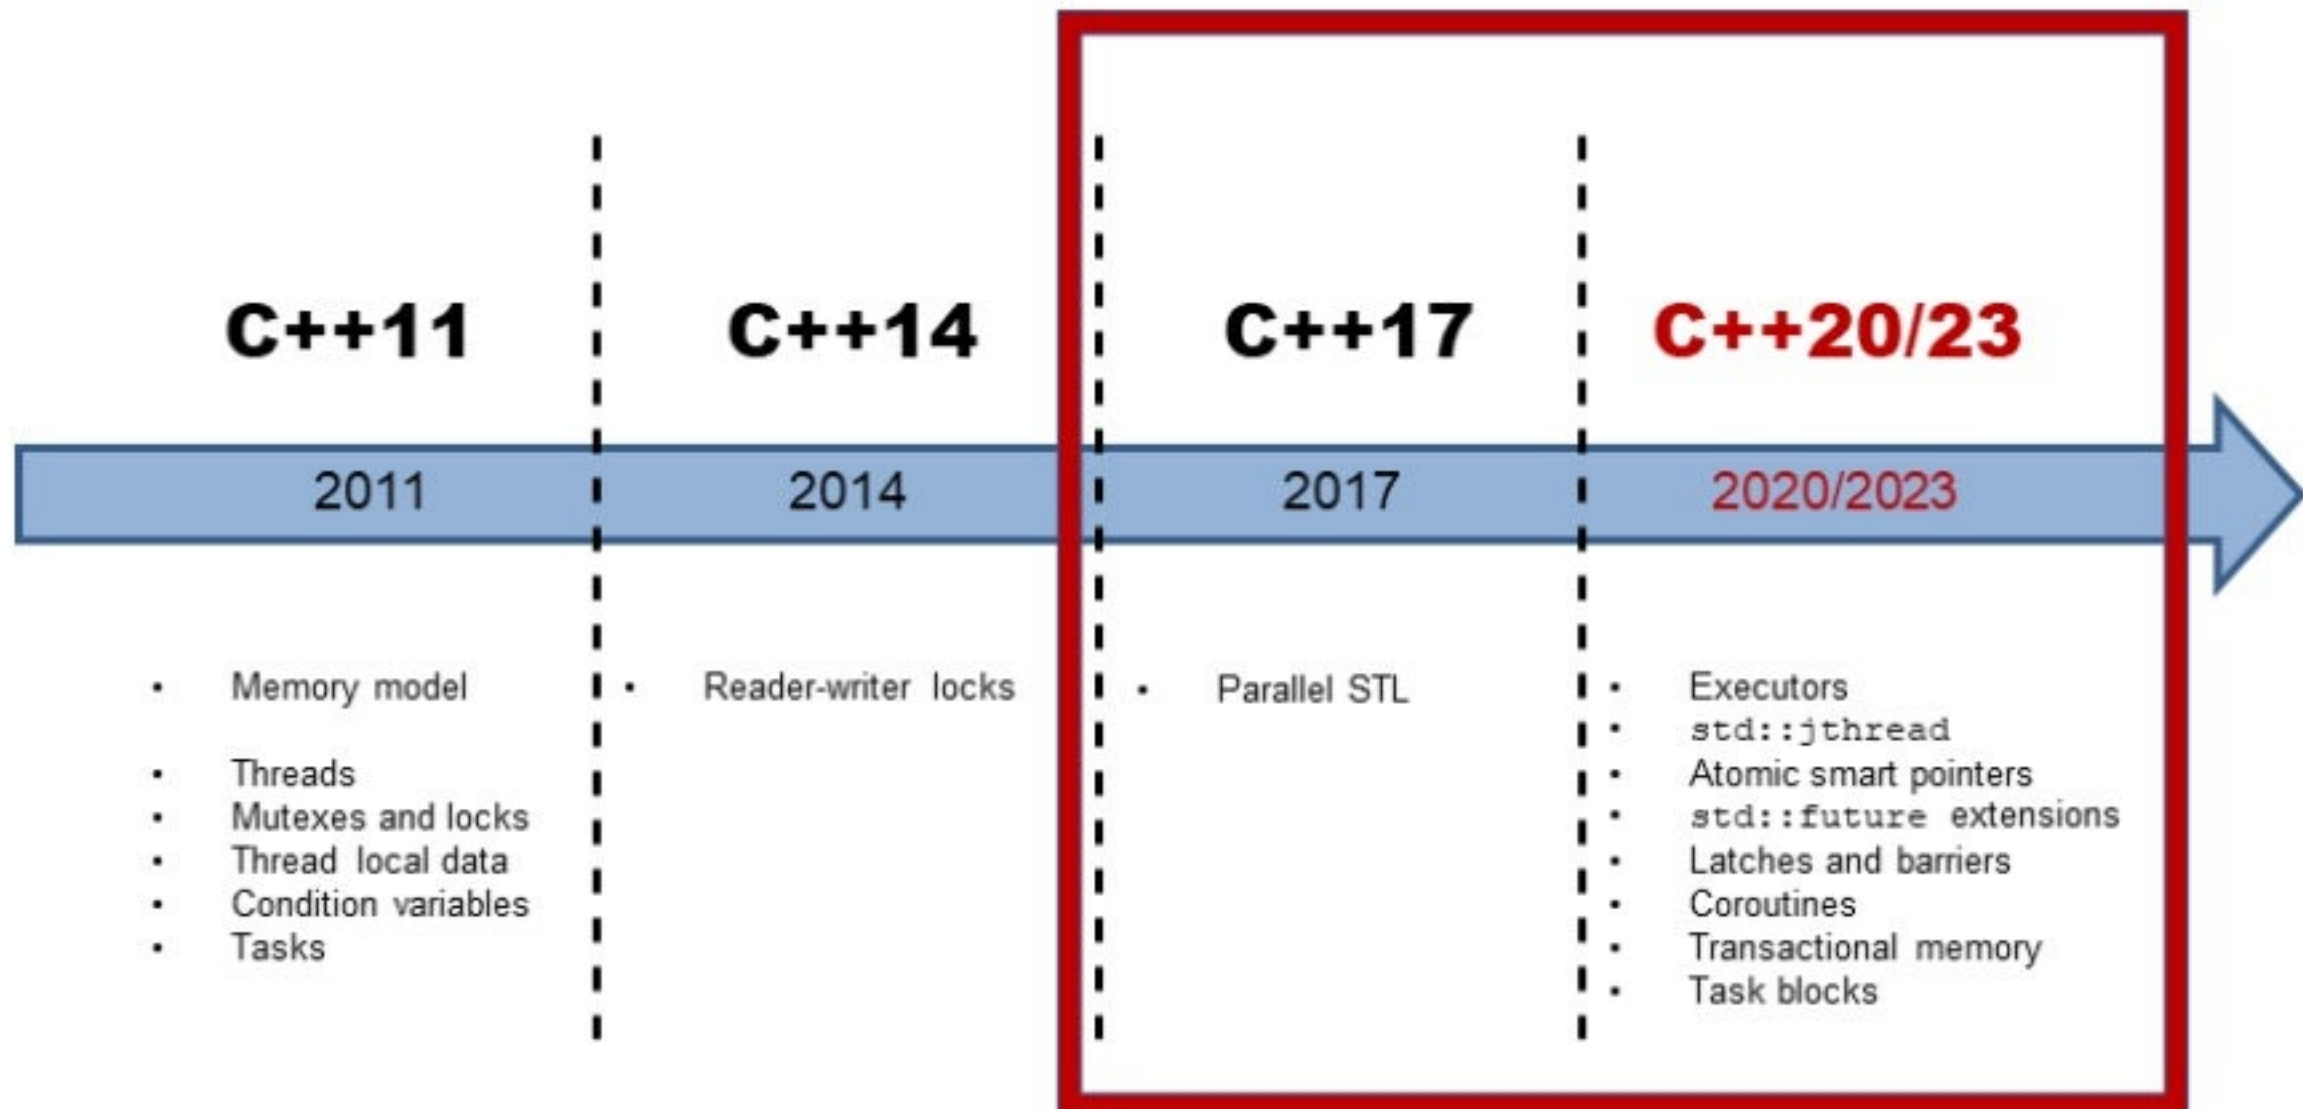
```
template <typename Func>
int traverse(node& n, Func && f){
    int left = 0, right = 0;
    define_task_block(
        [&](task_block& tb) {
            if (n.left) tb.run([&]{ left = traverse(*n.left, f); });
            if (n.right) tb.run([&]{ right = traverse(*n.right, f); });
        }
    );
    return f(n) + left + right;
}
```

- **define\_task\_block**
  - tasks can be performed
  - tasks will be synchronised at the end of the task block
- **run**: starts a task

# Concurrency and Parallelism in C++

## Multithreading

## Concurrency and Parallelism



# Concurrency and Parallelism in C++





# Proposals

- Atomic smart pointers: [N4162](#) (2014)
- `std::future` extensions: [N4107](#) (2014) and [P070r1](#) (2017)
- Latches and barriers: [P0666R0](#) (2017)
- Coroutines: [N4723](#) (2018)
- Transactional memory: [N4265](#) (2014)
- Task blocks: [N4411](#) (2015)
- Executors: [P0761](#) (2018)
- Concurrent unordered associative containers: [N3732](#) (2013)
- Concurrent Queue: [P0260r0](#) (2016)
- Pipelines: [N3534](#) (2013)
- Distributed counters: [P0261r1](#) (2016)
- Jthread [P0660R2](#) (2018)



# Blogs

[www.grimm-jaud.de](http://www.grimm-jaud.de) [De]

[www.ModernesCpp.com](http://www.ModernesCpp.com) [En]

Rainer Grimm

Training, Coaching, and  
Technology Consulting

[www.ModernesCpp.de](http://www.ModernesCpp.de)